

**Fine-tuning code** using **MobileNetV2** in **TensorFlow/Keras**, designed for a **binary image classification task** like **cats vs dogs**.

✓ **What this does:**

- Loads pre-trained MobileNetV2.
- Adds new layers for your task.
- Trains those new layers first.
- Then fine-tunes the entire model.

---

📦 **Requirements**

Install TensorFlow if you haven't:

```
pip install tensorflow
```

You'll also need a directory like:

```
data/
```

```
└─ cats_and_dogs/
```

```
    └─ cats/
```

```
        └─ cat1.jpg
```

```
        └─ ...
```

```
    └─ dogs/
```

```
        └─ dog1.jpg
```

```
        └─ ...
```

---

✓ **Full Code**

```
import tensorflow as tf
```

```
from tensorflow.keras.applications import MobileNetV2
```

```
from tensorflow.keras.models import Model
```

```
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
import os
```

```
# -----
```

```
# 1. Load Pretrained Base Model
```

```

# -----

base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
base_model.trainable = False # Freeze base layers initially

# -----

# 2. Add Custom Top Layers
# -----

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(128, activation='relu')(x)
output = Dense(1, activation='sigmoid')(x) # For binary classification

model = Model(inputs=base_model.input, outputs=output)

# -----

# 3. Compile Model (initial training)
# -----

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])

# -----

# 4. Prepare Data
# -----

data_dir = 'data/cats_and_dogs'

train_datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.2,
    horizontal_flip=True,
    zoom_range=0.2

```

```
)
```

```
train_generator = train_datagen.flow_from_directory(  
    data_dir,  
    target_size=(224, 224),  
    batch_size=32,  
    class_mode='binary',  
    subset='training'  
)
```

```
val_generator = train_datagen.flow_from_directory(  
    data_dir,  
    target_size=(224, 224),  
    batch_size=32,  
    class_mode='binary',  
    subset='validation'  
)
```

```
# -----
```

```
# 5. Train Top Layers Only
```

```
# -----
```

```
print("Training new classification head...")
```

```
model.fit(train_generator, epochs=5, validation_data=val_generator)
```

```
# -----
```

```
# 6. Fine-tune the Whole Model
```

```
# -----
```

```
base_model.trainable = True # Unfreeze the base model
```

```
# Set a lower learning rate for fine-tuning
```

```
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=1e-5),
```

```
        loss='binary_crossentropy',
        metrics=['accuracy'])

print("Fine-tuning entire model...")
model.fit(train_generator, epochs=5, validation_data=val_generator)

# -----
# 7. Save Model
# -----

model.save("fine_tuned_mobilenetv2.h5")
print("Model saved as fine_tuned_mobilenetv2.h5")
```

---



### Notes

- You can increase epochs as needed.
  - You can use `model.evaluate()` to test on unseen data.
  - If your dataset is not in folders, I can help you adapt this for CSV/image-label pairs.
-