

FORWARD AND BACKWARD PROPAGATION

✓ Improvements:

- Multiple training samples
 - Multiple training epochs
 - Loss printed each epoch
 - Visualizing loss over time (optional)
-

🧠 Problem Setup:

We'll use a very simple dataset (X, y) and keep:

- 2 inputs
- 1 hidden layer (2 neurons)
- 1 output neuron

📌 Full Code with Multiple Epochs:

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
# Activation function: Sigmoid
```

```
def sigmoid(x):
```

```
    return 1 / (1 + np.exp(-x))
```

```
# Derivative of sigmoid
```

```
def sigmoid_derivative(x):
```

```
    return x * (1 - x)
```

```
# Dataset: 4 training examples (XOR-like pattern)
```

```
X = np.array([
```

```
    [0.05, 0.10],
```

```
    [0.10, 0.20],
```

```
    [0.20, 0.05],
```

```
    [0.30, 0.25]
```

```
])
```

```
y_true = np.array([
    [0.01],
    [0.99],
    [0.01],
    [0.99]
])
```

```
# Initialize weights and biases with small random values
```

```
np.random.seed(42) # for reproducibility
```

```
W1 = np.random.rand(2, 2)
```

```
b1 = np.random.rand(1, 2)
```

```
W2 = np.random.rand(2, 1)
```

```
b2 = np.random.rand(1, 1)
```

```
# Hyperparameters
```

```
lr = 0.5
```

```
epochs = 10000
```

```
loss_history = []
```

```
# Training loop
```

```
for epoch in range(epochs):
```

```
    # ---- FORWARD PROPAGATION ----
```

```
    hidden_input = np.dot(X, W1) + b1
```

```
    hidden_output = sigmoid(hidden_input)
```

```
    final_input = np.dot(hidden_output, W2) + b2
```

```
    final_output = sigmoid(final_input)
```

```

# ---- LOSS ----

loss = np.mean(0.5 * (y_true - final_output) ** 2)
loss_history.append(loss)

# ---- BACKPROPAGATION ----

error_output = final_output - y_true
delta_output = error_output * sigmoid_derivative(final_output)

dW2 = np.dot(hidden_output.T, delta_output)
db2 = np.sum(delta_output, axis=0, keepdims=True)

error_hidden = np.dot(delta_output, W2.T)
delta_hidden = error_hidden * sigmoid_derivative(hidden_output)

dW1 = np.dot(X.T, delta_hidden)
db1 = np.sum(delta_hidden, axis=0, keepdims=True)

# ---- UPDATE WEIGHTS ----

W2 -= lr * dW2
b2 -= lr * db2
W1 -= lr * dW1
b1 -= lr * db1

# Print loss occasionally

if epoch % 1000 == 0:
    print(f"Epoch {epoch}, Loss: {loss:.6f}")

# Final predictions
print("\nFinal predictions after training:")
print(final_output)

```

```
# Plot loss curve

plt.plot(loss_history)

plt.title("Loss over Epochs")

plt.xlabel("Epoch")

plt.ylabel("Loss")

plt.grid(True)

plt.show()
```

Output:

- Prints the loss every 1000 epochs
 - Displays a loss curve showing how training improves
 - Shows final output after training (should be close to targets)
-

What You've Learned:

- How to implement forward and backward propagation from scratch
 - How to train on multiple data samples
 - How to observe loss decreasing over time
 - How to make predictions from a trained model
-